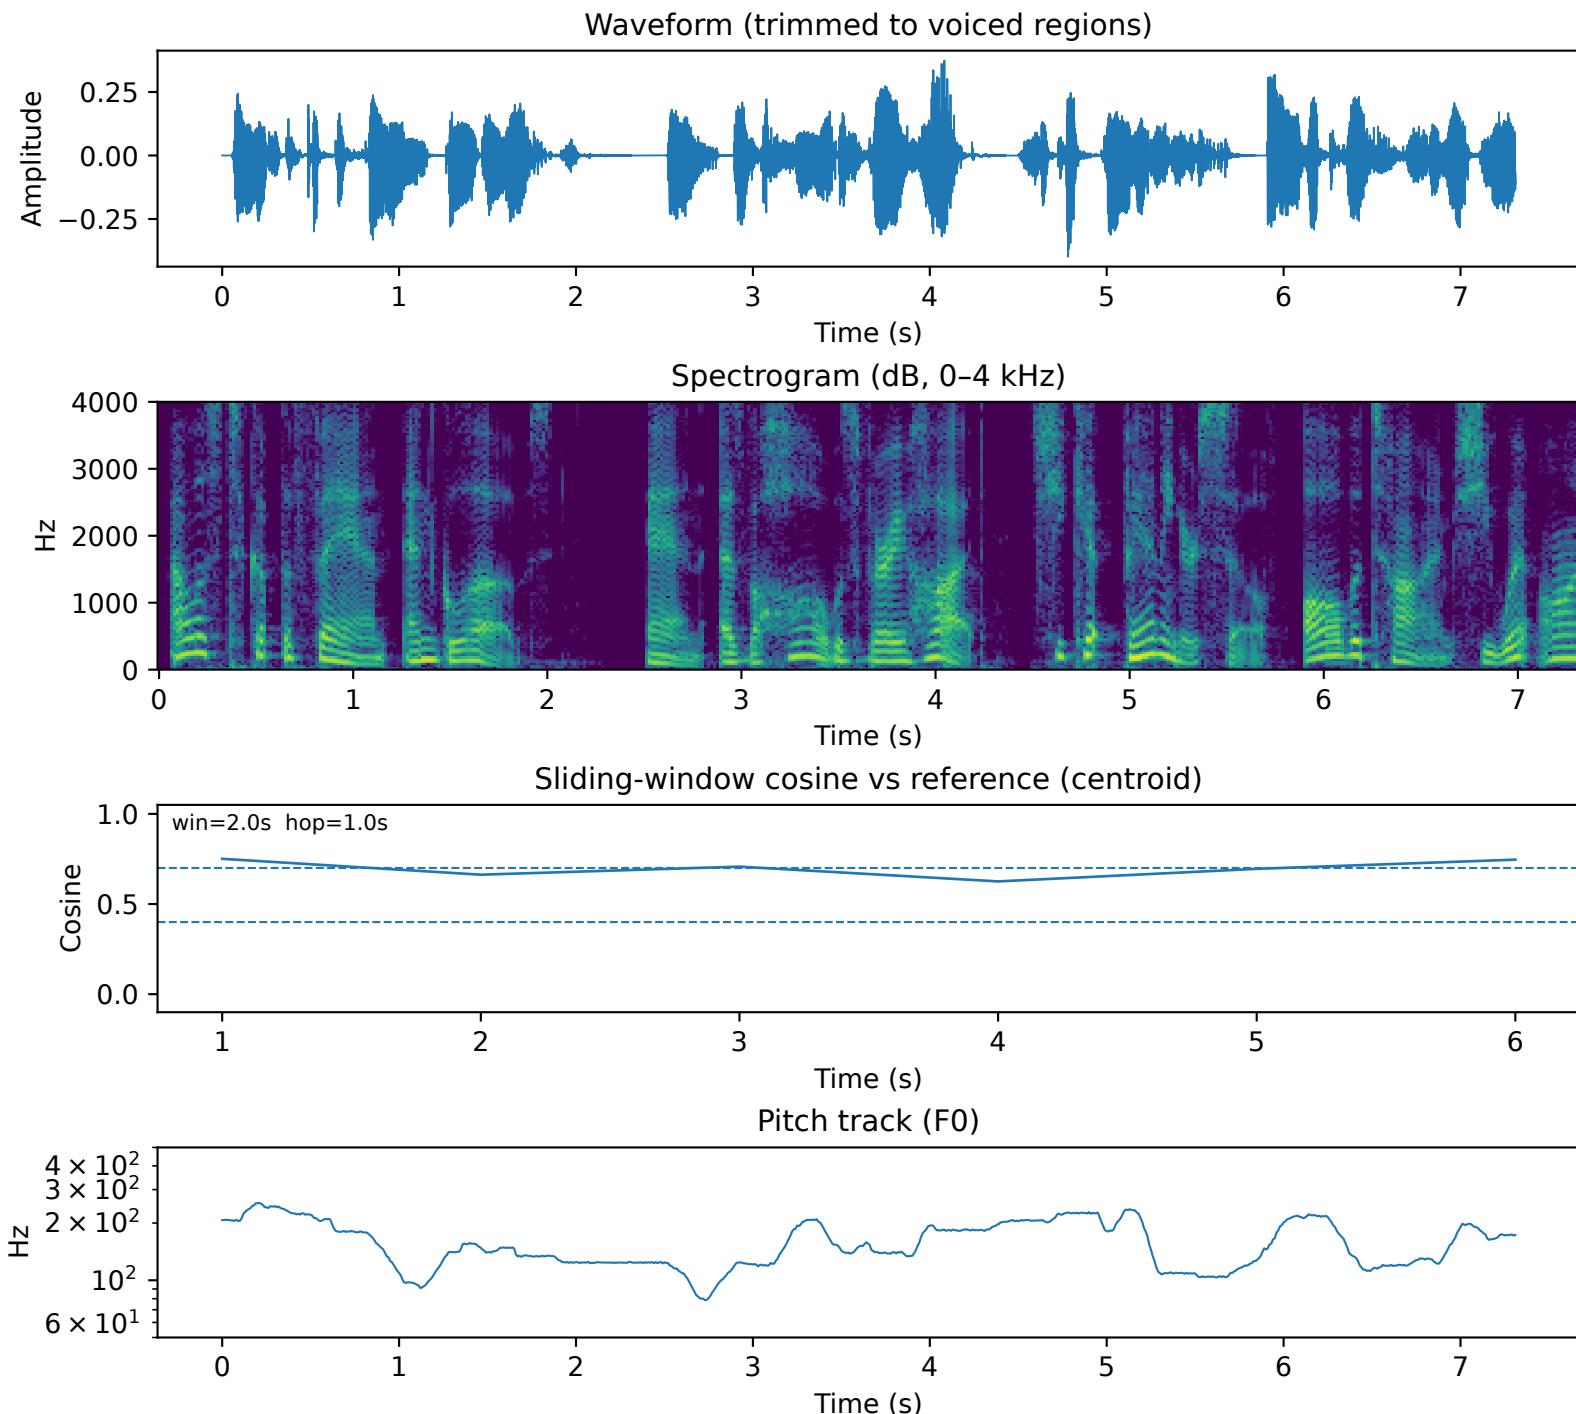


# Forensic Voice Report — label: libris\_260\_clean

Report ID: 20251027T131347Z\_dc3df386 | File: 260-123440-0004.wav | SHA256: c70586705576...  
Voiced: 7.31s (VR=0.94) | SNR: 8.5 dB | Scores - Best: 0.94 | Top3 mean: 0.52 | Centroid: 0.90  
Decision: high\_match (thr: med=0.40, high=0.70) | CM: 0.34 | F0 median: 147.27 Hz (IQR 72.10 Hz)



## Per-segment (voiced $\geq 0.2$ s)

| Start (s) | End (s) | Dur (s) | Cosine | Flag |
|-----------|---------|---------|--------|------|
| 0.00      | 7.00    | 7.00    | 0.70   | HIGH |

See next page for Top Enrollments and Full Provenance / Methods.

# Top Enrollments and Provenance

## Top enrollments by cosine

| Enroll ID | Created             | Cosine |
|-----------|---------------------|--------|
| a5f19f64  | 2025-10-25 23:15:04 | 0.94   |
| d485c61c  | 2025-10-25 23:15:12 | 0.32   |
| 1b32d906  | 2025-10-25 23:06:20 | 0.29   |
| 43c245c5  | 2025-10-25 23:14:51 | 0.29   |
| bc4bac39  | 2025-10-25 23:32:58 | 0.24   |
| f8645e67  | 2025-10-25 23:34:09 | 0.20   |

## Provenance & Methods

```
File: 260-123440-0004.wav Bytes: 250638 | SHA256:  
c70586705576bd458accabf9addfb4bf6cac497816a85e3e2fc3f54f984d16b6 Orig sample rate: 16000 Hz | Channels: 1 |  
Bit depth: 16 Duration: 7.83 s System: Windows 11 | Host: DESKTOP-32VU5AV | User: asus Python: 3.13.7 |  
numpy: 2.3.3 | torch: 2.8.0+cpu | torchaudio: 2.8.0+cpu speechbrain: 1.0.3 | torchcrepe: None Config  
snapshot: thresholds (med=0.40, high=0.70); VAD mode=2, pad_ms=50; scoring: best & mean(top-3) vs enrollments;  
centroid cosine.
```

```
Speech isolation (VAD): WebRTC VAD, 20 ms frames, mode 2, silence trimmed with ±50 ms padding. Voiced ratio =  
fraction of voiced frames. Embedding extraction: ECAPA (speechbrain/spkrec-ecapa-voxceleb) at 16 kHz mono;  
embeddings L2-normalized. Similarity scoring: primary = cosine to each enrollment for the label; report best  
and mean(top-3). Secondary = cosine to label centroid (mean of enrollments). Decision policy: thresholds  
medium=0.40, high=0.70 → label as low/medium/high match. Anti-spoofing (optional): CM score is reported if  
available.
```